

heart-failure-prediction

Use the "Run" button to execute the code.

```
!pip install jovian --upgrade --quiet
```

```
import jovian
```

```
# Execute this to save new versions of the notebook
jovian.commit(project="heart-failure-prediction")
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
!pip install opendatasets
```

Collecting opendatasets

Using cached opendatasets-0.1.20-py3-none-any.whl (14 kB)

Collecting kaggle

Downloading kaggle-1.5.12.tar.gz (58 kB)

58 kB 4.6 MB/s eta 0:00:011

Requirement already satisfied: click in /opt/conda/lib/python3.9/site-packages (from opendatasets) (8.0.1)

Requirement already satisfied: tqdm in /opt/conda/lib/python3.9/site-packages (from opendatasets) (4.62.3)

```
Requirement already satisfied: six>=1.10 in /opt/conda/lib/python3.9/site-packages (from
kaggle->opendatasets) (1.16.0)
```

Requirement already satisfied: certifi in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2021.5.30)

Requirement already satisfied: python-dateutil in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2.8.2)

Requirement already satisfied: requests in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (2.26.0)

Collecting python-slugify

Downloading python_slugify-5.0.2-py2.py3-none-any.whl (6.7 kB)

Requirement already satisfied: urllib3 in /opt/conda/lib/python3.9/site-packages (from kaggle->opendatasets) (1.26.7)

Collecting text-unidecode>=1.3

Downloading text_unidecode-1.3-py2.py3-none-any.whl (78 kB)

```
|██████████| 78 kB 10.2 MB/s eta 0:00:01
```

Requirement already satisfied: charset-normalizer~=2.0.0 in /opt/conda/lib/python3.9/site-packages (from requests->kaggle->opendatasets) (2.0.0)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.9/site-packages (from requests->kaggle->opendatasets) (3.1)
Building wheels for collected packages: kaggle
Building wheel for kaggle (setup.py) ... done
Created wheel for kaggle: filename=kaggle-1.5.12-py3-none-any.whl size=73051 sha256=b19c542fbd70481d0251e09501d6954175725b422dffb9ebc3a632aa55d57864
Stored in directory:
/home/jovyan/.cache/pip/wheels/ac/b2/c3/fa4706d469b5879105991d1c8be9a3c2ef329ba9fe2ce506
Successfully built kaggle
Installing collected packages: text-unidecode, python-slugify, kaggle, opendatasets
Successfully installed kaggle-1.5.12 opendatasets-0.1.20 python-slugify-5.0.2 text-unidecode-1.3

```
import opendatasets as od
```

```
od.download('https://www.kaggle.com/fedesoriano/heart-failure-prediction')
```

Please provide your Kaggle credentials to download this dataset. Learn more:

<http://bit.ly/kaggle-creds>

Your Kaggle username: swetsheers

Your Kaggle Key:

Downloading heart-failure-prediction.zip to ./heart-failure-prediction

100%|██████████| 8.56k/8.56k [00:00<00:00, 5.85MB/s]

```
import os
import pandas as pd
os.listdir('./heart-failure-prediction')
```

```
['heart.csv']
```

```
heart=pd.read_csv('./heart-failure-prediction/heart.csv')
```

```
heart
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
0	40	M	ATA	140	289	0	Normal	172	N	0.0	
1	49	F	NAP	160	180	0	Normal	156	N	1.0	
2	37	M	ATA	130	283	0	ST	98	N	0.0	
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
	4	54	M	NAP	150	195	0	Normal	122	N	0.0

	913	45	M	TA	110	264	0	Normal	132	N	1.2
	914	68	M	ASY	144	193	1	Normal	141	N	3.4
	915	57	M	ASY	130	131	0	Normal	115	Y	1.2
	916	57	F	ATA	130	236	0	LVH	174	N	0.0
	917	38	M	NAP	138	175	0	Normal	173	N	0.0

918 rows × 12 columns

```
heart.isna().sum()
```

```
Age      0
Sex      0
ChestPainType  0
RestingBP      0
Cholesterol    0
FastingBS     0
RestingECG     0
MaxHR         0
ExerciseAngina 0
Oldpeak       0
ST_Slope      0
HeartDisease   0
dtype: int64
```

```
heart.describe()
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease
count	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000
mean	53.510893	132.396514	198.799564	0.233115	136.809368	0.887364	0.553377
std	9.432617	18.514154	109.384145	0.423046	25.460334	1.066570	0.497414
min	28.000000	0.000000	0.000000	0.000000	60.000000	-2.600000	0.000000
25%	47.000000	120.000000	173.250000	0.000000	120.000000	0.000000	0.000000
50%	54.000000	130.000000	223.000000	0.000000	138.000000	0.600000	1.000000
75%	60.000000	140.000000	267.000000	0.000000	156.000000	1.500000	1.000000
max	77.000000	200.000000	603.000000	1.000000	202.000000	6.200000	1.000000

```
heart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 918 entries, 0 to 917
```

```
Data columns (total 12 columns):
```

```
#    Column          Non-Null Count  Dtype
```

---	-----	-----	-----
0	Age	918 non-null	int64
1	Sex	918 non-null	object
2	ChestPainType	918 non-null	object
3	RestingBP	918 non-null	int64
4	Cholesterol	918 non-null	int64
5	FastingBS	918 non-null	int64
6	RestingECG	918 non-null	object
7	MaxHR	918 non-null	int64
8	ExerciseAngina	918 non-null	object
9	Oldpeak	918 non-null	float64
10	ST_Slope	918 non-null	object
11	HeartDisease	918 non-null	int64

dtypes: float64(1), int64(6), object(5)

memory usage: 86.2+ KB

```
jovian.commit()
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib
import os
%matplotlib inline

pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 150)
sns.set_style('darkgrid')
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (10, 6)
matplotlib.rcParams['figure.facecolor'] = '#00000000'
```

```
!pip install plotly
```

Collecting plotly

Using cached plotly-5.5.0-py2.py3-none-any.whl (26.5 MB)

Collecting tenacity>=6.2.0

Downloading tenacity-8.0.1-py3-none-any.whl (24 kB)

Requirement already satisfied: six in /opt/conda/lib/python3.9/site-packages (from

plotly) (1.16.0)

Installing collected packages: tenacity, plotly

Successfully installed plotly-5.5.0 tenacity-8.0.1


```
import plotly.express as px
```

```
sns.scatterplot(heart.Age,  
                heart.Cholesterol,  
                hue=heart.HeartDisease  
                );
```

```
# Chart title  
plt.title("Heart");
```

/opt/conda/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning:

Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



```
px.scatter(heart,x='Age',y='Cholesterol',color='HeartDisease')
```

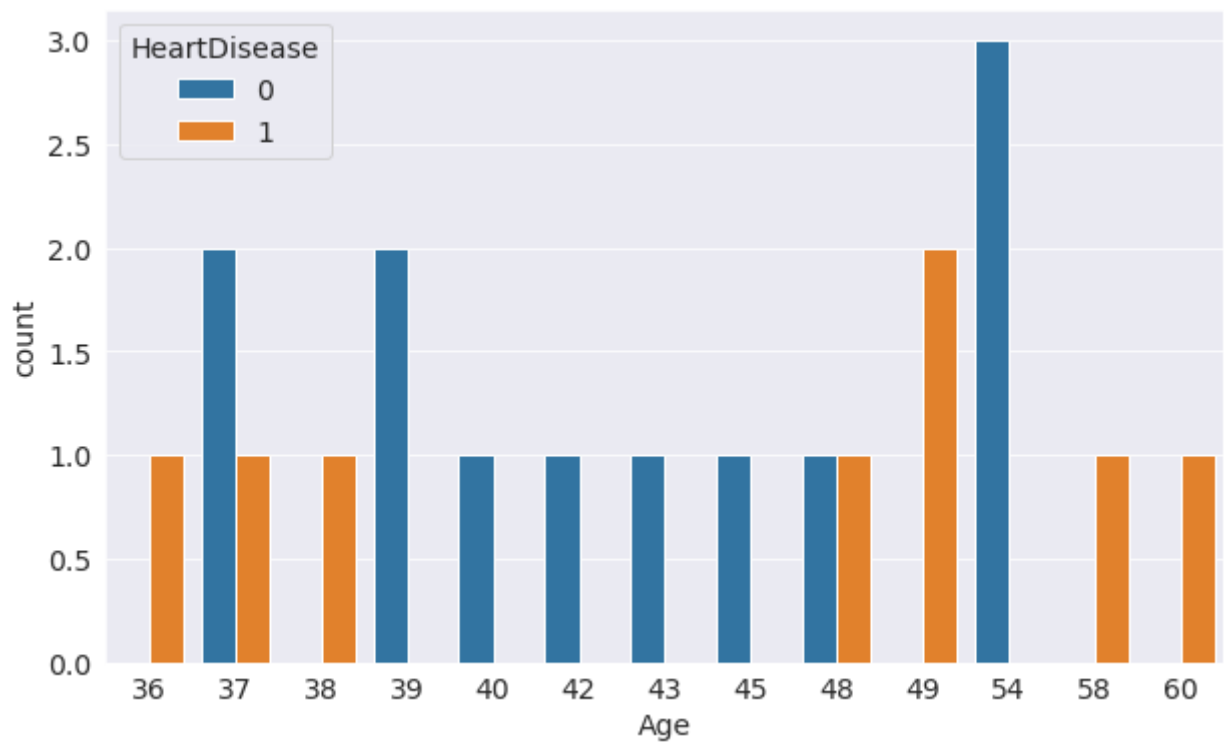
heart

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
0	40	M	ATA	140	289	0	Normal	172	N	0.0	
1	49	F	NAP	160	180	0	Normal	156	N	1.0	
2	37	M	ATA	130	283	0	ST	98	N	0.0	
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	
4	54	M	NAP	150	195	0	Normal	122	N	0.0	
...	
913	45	M	TA	110	264	0	Normal	132	N	1.2	
914	68	M	ASY	144	193	1	Normal	141	N	3.4	
915	57	M	ASY	130	131	0	Normal	115	Y	1.2	
916	57	F	ATA	130	236	0	LVH	174	N	0.0	
917	38	M	NAP	138	175	0	Normal	173	N	0.0	

918 rows × 12 columns

```
sns.countplot(x=heart['Age'].head(20),hue=heart['HeartDisease'])
```

<AxesSubplot:xlabel='Age', ylabel='count'>



```
heart.columns
```

```
Index(['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol', 'FastingBS',
      'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak', 'ST_Slope',
      'HeartDisease'],
      dtype='object')
```

```
input_cols=['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol', 'FastingBS',
            'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak', 'ST_Slope']
```

```
target='HeartDisease'
```

```
x_train=heart[input_cols]
target=heart[target]
```

```
x_train
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
0	40	M	ATA	140	289	0	Normal	172	N	0.0	
1	49	F	NAP	160	180	0	Normal	156	N	1.0	
2	37	M	ATA	130	283	0	ST	98	N	0.0	
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	
4	54	M	NAP	150	195	0	Normal	122	N	0.0	
...	
913	45	M	TA	110	264	0	Normal	132	N	1.2	
914	68	M	ASY	144	193	1	Normal	141	N	3.4	
915	57	M	ASY	130	131	0	Normal	115	Y	1.2	

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
916	57	F	ATA	130	236	0	LVH	174	N	0.0	
917	38	M	NAP	138	175	0	Normal	173	N	0.0	

918 rows × 11 columns

```
heart.nunique()
```

```
Age          50
Sex           2
ChestPainType 4
RestingBP     67
Cholesterol   222
FastingBS     2
RestingECG    3
MaxHR        119
ExerciseAngina 2
Oldpeak       53
ST_Slope      3
HeartDisease  2
dtype: int64
```

```
numeric_cols=['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']
cat_cols=['Sex', 'ChestPainType', 'FastingBS', 'RestingECG', 'ExerciseAngina', 'ST_Slope']
```

```
!pip install sklearn
from sklearn.impute import SimpleImputer
```

```
Requirement already satisfied: sklearn in /opt/conda/lib/python3.9/site-packages (0.0)
Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.9/site-packages
(from sklearn) (1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.9/site-
packages (from scikit-learn->sklearn) (2.2.0)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.9/site-packages
(from scikit-learn->sklearn) (1.0.1)
Requirement already satisfied: numpy>=1.14.6 in /opt/conda/lib/python3.9/site-packages
(from scikit-learn->sklearn) (1.20.3)
Requirement already satisfied: scipy>=1.1.0 in /opt/conda/lib/python3.9/site-packages
(from scikit-learn->sklearn) (1.7.1)
```

```
imputer=SimpleImputer(strategy='mean')
```

```
imputer.fit(heart[numeric_cols])
```

```
SimpleImputer()
```



```
heart[numeric_cols]=imputer.transform(x_train[numeric_cols])
```

```
x_train[numeric_cols]
```

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak
0	40	140	289	172	0.0
1	49	160	180	156	1.0
2	37	130	283	98	0.0
3	48	138	214	108	1.5
4	54	150	195	122	0.0
...
913	45	110	264	132	1.2
914	68	144	193	141	3.4
915	57	130	131	115	1.2
916	57	130	236	174	0.0
917	38	138	175	173	0.0

918 rows × 5 columns

```
from sklearn.preprocessing import OneHotEncoder
```

```
encoder=OneHotEncoder(sparse=False, handle_unknown='ignore').fit(x_train[cat_cols])
```

```
encoded_cols=list(encoder.get_feature_names(cat_cols))
```

/opt/conda/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning:

Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please use get_feature_names_out instead.

```
len(encoded_cols)
```

16

```
x_train[encoded_cols]=encoder.transform(x_train[cat_cols])
```

```
x_train
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST
0	40	M	ATA	140	289	0	Normal	172	N	0.0	
1	49	F	NAP	160	180	0	Normal	156	N	1.0	

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
913	0.346939	0.55	0.437811	0.507042	0.431818	0.0	1.0	0.0	0.0
914	0.816327	0.72	0.320066	0.570423	0.681818	0.0	1.0	1.0	0.0
915	0.591837	0.65	0.217247	0.387324	0.431818	0.0	1.0	1.0	0.0
916	0.591837	0.65	0.391376	0.802817	0.295455	1.0	0.0	0.0	1.0
917	0.204082	0.69	0.290216	0.795775	0.295455	0.0	1.0	0.0	0.0

918 rows × 21 columns

```
jovian.commit()
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
from sklearn.model_selection import train_test_split
```

```
train_df, val_df, train_target, val_target = train_test_split(x_train, target, test_size=0.
```

```
train_df
```

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
557	0.571429	0.685	0.344942	0.436620	0.500000	0.0	1.0	0.0	0.0
260	0.367347	0.700	0.456053	0.739437	0.295455	0.0	1.0	0.0	1.0
235	0.224490	0.600	0.331675	0.704225	0.409091	0.0	1.0	0.0	1.0
218	0.551020	0.700	0.325041	0.633803	0.295455	0.0	1.0	0.0	1.0
382	0.306122	0.575	0.000000	0.598592	0.522727	0.0	1.0	1.0	0.0
...
106	0.408163	0.600	0.421227	0.352113	0.295455	1.0	0.0	1.0	0.0
270	0.346939	0.600	0.373134	0.563380	0.295455	0.0	1.0	1.0	0.0
860	0.653061	0.650	0.419569	0.591549	0.454545	0.0	1.0	1.0	0.0
435	0.653061	0.760	0.000000	0.408451	0.295455	0.0	1.0	1.0	0.0
102	0.244898	0.750	0.650083	0.492958	0.522727	1.0	0.0	1.0	0.0

826 rows × 21 columns

```
val_df
```

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
668	0.714286	0.700	0.323383	0.838028	0.295455	1.0	0.0	0.0	1.0
30	0.510204	0.725	0.859038	0.492958	0.295455	0.0	1.0	0.0	0.0

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
377	0.755102	0.800	0.000000	0.436620	0.431818	0.0	1.0	1.0	0.0
535	0.571429	0.650	0.000000	0.436620	0.409091	0.0	1.0	1.0	0.0
807	0.530612	0.540	0.512438	0.676056	0.295455	0.0	1.0	0.0	1.0
793	0.795918	0.625	0.421227	0.725352	0.318182	0.0	1.0	1.0	0.0
363	0.571429	0.600	0.000000	0.619718	0.295455	0.0	1.0	1.0	0.0
583	0.836735	0.710	0.449420	0.464789	0.329545	0.0	1.0	0.0	0.0
165	0.367347	0.700	0.451078	0.809859	0.522727	0.0	1.0	0.0	0.0
483	0.612245	0.600	0.000000	0.323944	0.465909	0.0	1.0	1.0	0.0
773	0.571429	0.600	0.320066	0.718310	0.511364	0.0	1.0	0.0	0.0
551	0.693878	0.600	0.364842	0.183099	0.295455	0.0	1.0	0.0	0.0
768	0.734694	0.650	0.502488	0.436620	0.522727	1.0	0.0	1.0	0.0
694	0.571429	0.600	0.391376	0.830986	0.386364	0.0	1.0	0.0	1.0
718	0.591837	0.825	0.479270	0.450704	0.409091	0.0	1.0	1.0	0.0
312	0.265306	0.625	0.000000	0.816901	0.477273	0.0	1.0	1.0	0.0
713	0.734694	0.700	0.519071	0.514085	0.318182	1.0	0.0	0.0	0.0
309	0.591837	0.475	0.000000	0.859155	0.375000	0.0	1.0	1.0	0.0
846	0.224490	0.590	0.363184	0.563380	0.431818	0.0	1.0	1.0	0.0
616	0.795918	0.575	0.935323	0.704225	0.477273	1.0	0.0	0.0	0.0
355	0.795918	0.725	0.000000	0.457746	0.295455	0.0	1.0	0.0	0.0
39	0.408163	0.750	0.376451	0.492958	0.409091	1.0	0.0	1.0	0.0
231	0.244898	0.650	0.466003	0.753521	0.295455	0.0	1.0	0.0	0.0
822	0.612245	0.525	0.398010	0.661972	0.363636	0.0	1.0	0.0	0.0
603	0.938776	0.775	0.514096	0.366197	0.465909	0.0	1.0	1.0	0.0
63	0.367347	0.600	0.459370	0.457746	0.409091	0.0	1.0	1.0	0.0
192	0.408163	0.650	0.406302	0.704225	0.295455	0.0	1.0	0.0	1.0
481	0.836735	0.700	0.000000	0.408451	0.579545	0.0	1.0	0.0	0.0
866	0.326531	0.650	0.363184	0.901408	0.295455	0.0	1.0	0.0	1.0
67	0.081633	0.550	0.373134	0.873239	0.295455	0.0	1.0	0.0	1.0
72	0.489796	0.600	0.301824	0.633803	0.295455	0.0	1.0	1.0	0.0
655	0.244898	0.760	0.369818	0.852113	0.295455	0.0	1.0	1.0	0.0
679	0.714286	0.725	0.386401	0.633803	0.556818	0.0	1.0	0.0	0.0
139	0.306122	0.750	0.409619	0.492958	0.522727	0.0	1.0	1.0	0.0
732	0.571429	1.000	0.477612	0.514085	0.750000	1.0	0.0	1.0	0.0
824	0.183673	0.650	0.414594	0.894366	0.693182	0.0	1.0	0.0	0.0
174	0.489796	0.700	0.441128	0.521127	0.522727	0.0	1.0	1.0	0.0
896	0.387755	0.650	0.419569	0.838028	0.295455	0.0	1.0	0.0	0.0
499	0.693878	0.675	0.492537	0.492958	0.409091	0.0	1.0	1.0	0.0
70	0.591837	0.700	0.439469	0.598592	0.409091	0.0	1.0	0.0	1.0
716	0.795918	0.600	0.393035	0.077465	0.409091	0.0	1.0	1.0	0.0

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
23	0.326531	0.750	0.477612	0.633803	0.636364	0.0	1.0	0.0	1.0
541	0.979592	0.520	0.187396	0.422535	0.693182	0.0	1.0	0.0	0.0
799	0.510204	0.650	0.407960	0.795775	0.295455	0.0	1.0	0.0	0.0
672	0.653061	0.600	0.295191	0.253521	0.295455	1.0	0.0	0.0	0.0
826	0.469388	0.625	0.406302	0.746479	0.568182	0.0	1.0	0.0	0.0
250	0.326531	0.675	0.814262	0.528169	0.295455	0.0	1.0	1.0	0.0
752	0.571429	0.625	0.412935	0.591549	0.431818	0.0	1.0	1.0	0.0
350	0.510204	0.600	0.000000	0.422535	0.295455	0.0	1.0	1.0	0.0
758	0.469388	0.625	0.353234	0.457746	0.454545	0.0	1.0	0.0	0.0
759	0.530612	0.960	0.469320	0.950704	0.295455	0.0	1.0	0.0	1.0
107	0.122449	0.750	0.354892	0.760563	0.295455	0.0	1.0	0.0	1.0
445	0.551020	0.680	0.378109	0.450704	0.477273	0.0	1.0	0.0	0.0
141	0.448980	0.700	0.565506	0.457746	0.579545	0.0	1.0	1.0	0.0
650	0.408163	0.650	0.424544	0.633803	0.295455	0.0	1.0	1.0	0.0
544	0.673469	0.700	0.494196	0.422535	0.295455	1.0	0.0	0.0	1.0
110	0.632653	0.650	0.311774	0.450704	0.409091	1.0	0.0	0.0	1.0
593	0.734694	0.650	0.427861	0.492958	0.295455	0.0	1.0	1.0	0.0
519	0.714286	0.480	0.505804	0.429577	0.409091	0.0	1.0	1.0	0.0
907	0.326531	0.600	0.280265	0.591549	0.613636	0.0	1.0	1.0	0.0
675	0.591837	0.750	0.208955	0.795775	0.318182	0.0	1.0	0.0	0.0
280	0.653061	0.600	0.407960	0.528169	0.295455	0.0	1.0	0.0	0.0
136	0.306122	0.600	0.356551	0.809859	0.295455	1.0	0.0	0.0	1.0
422	0.755102	0.750	0.391376	0.316901	0.295455	0.0	1.0	1.0	0.0
208	0.000000	0.650	0.218905	0.880282	0.295455	0.0	1.0	0.0	1.0
442	0.469388	0.640	0.000000	0.457746	0.431818	0.0	1.0	1.0	0.0
86	0.755102	0.850	0.436153	0.366197	0.522727	0.0	1.0	1.0	0.0
44	0.306122	0.600	0.290216	0.422535	0.409091	0.0	1.0	1.0	0.0
531	0.734694	0.715	0.507463	0.387324	0.500000	0.0	1.0	1.0	0.0
913	0.346939	0.550	0.437811	0.507042	0.431818	0.0	1.0	0.0	0.0
634	0.244898	0.700	0.330017	0.830986	0.454545	0.0	1.0	0.0	0.0
290	0.408163	0.550	0.349917	0.549296	0.295455	0.0	1.0	0.0	0.0
338	0.714286	0.700	0.000000	0.626761	0.522727	0.0	1.0	1.0	0.0
357	0.510204	0.600	0.000000	0.246479	0.295455	0.0	1.0	0.0	1.0
292	0.510204	0.650	0.301824	0.619718	0.295455	0.0	1.0	1.0	0.0
227	0.204082	0.460	0.194030	0.521127	0.579545	0.0	1.0	1.0	0.0
591	0.612245	0.500	0.353234	0.352113	0.295455	0.0	1.0	1.0	0.0
425	0.653061	0.800	0.442786	0.683099	0.352273	0.0	1.0	0.0	1.0
789	0.122449	0.590	0.301824	0.802817	0.295455	0.0	1.0	0.0	0.0
522	0.448980	0.720	0.578773	0.422535	0.409091	0.0	1.0	1.0	0.0

	Age	RestingBP	Cholesterol	MaxHR	Oldpeak	Sex_F	Sex_M	ChestPainType_ASY	ChestPainType_ATA
861	0.755102	0.550	0.411277	0.690141	0.363636	0.0	1.0	1.0	0.0
352	0.571429	0.600	0.000000	0.281690	0.181818	0.0	1.0	1.0	0.0
493	0.469388	0.685	0.562189	0.471831	0.488636	0.0	1.0	0.0	0.0
60	0.428571	0.500	0.419569	0.802817	0.295455	0.0	1.0	0.0	1.0
598	0.551020	0.600	0.374793	0.471831	0.488636	0.0	1.0	1.0	0.0
722	0.653061	0.750	0.427861	0.683099	0.590909	1.0	0.0	1.0	0.0
426	0.571429	0.630	0.275290	0.563380	0.295455	0.0	1.0	0.0	1.0
468	0.693878	0.760	0.253731	0.260563	0.477273	0.0	1.0	1.0	0.0
66	0.346939	0.660	0.492537	0.591549	0.295455	1.0	0.0	1.0	0.0
332	0.204082	0.500	0.000000	0.838028	0.170455	0.0	1.0	0.0	0.0
375	0.918367	0.800	0.000000	0.429577	0.295455	1.0	0.0	0.0	0.0
381	0.448980	0.575	0.000000	0.422535	0.352273	0.0	1.0	1.0	0.0

```
train_target
```

```
557    1
260    0
235    0
218    0
382    1
..
106    0
270    0
860    1
435    0
102    1
```

Name: HeartDisease, Length: 826, dtype: int64

```
len(val_target)
```

92

```
jovian.commit()
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
from sklearn.linear_model import LogisticRegression
```

```
model=LogisticRegression(solver='liblinear')
```

```
model.fit(train_df, train_target)
```

```
LogisticRegression(solver='liblinear')
```

```
pred=model.predict(train_df)
```

```
print(model.coef_.tolist())
```

```
[[0.6636763929634307, -0.03624671646388043, -1.7420098735456697, -0.5594161008880127,
1.8043899424890677, -0.6827890695652646, 0.6041332315396141, 1.0777271131154214,
-0.7442328782759892, -0.4235185174417452, 0.011368444576650361, -0.6520137780667163,
0.5733579400410825, 0.16209948518749306, -0.04740677158267605, -0.19334855163046916,
-0.5489776117174245, 0.4703217736917812, -0.021606613743019556, 1.1714206628929604,
-1.2284698871756172]]
```

```
print(model.intercept_)
```

```
[-0.07865584]
```

```
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
accuracy_score(pred, train_target)
```

```
0.8619854721549637
```

```
pred1=model.predict(val_df)
```

```
accuracy_score(pred1, val_target)
```

```
0.8804347826086957
```

```
val_probs = model.predict_proba(val_df)
val_probs
```

```
array([[0.98258563, 0.01741437],
       [0.71131789, 0.28868211],
       [0.02626967, 0.97373033],
       [0.02647756, 0.97352244],
       [0.95692554, 0.04307446],
       [0.06346221, 0.93653779],
       [0.1273922 , 0.8726078 ],
       [0.88934738, 0.11065262],
       [0.16686088, 0.83313912],
       [0.06874992, 0.93125008],
       [0.27087716, 0.72912284],
       [0.8734168 , 0.1265832 ],
       [0.37122695, 0.62877305],
```

[0.94200347, 0.05799653],
[0.04856241, 0.95143759],
[0.28713702, 0.71286298],
[0.97848585, 0.02151415],
[0.10619948, 0.89380052],
[0.18464867, 0.81535133],
[0.84646701, 0.15353299],
[0.19025218, 0.80974782],
[0.21266814, 0.78733186],
[0.94960411, 0.05039589],
[0.22555601, 0.77444399],
[0.15607651, 0.84392349],
[0.07947952, 0.92052048],
[0.95333256, 0.04666744],
[0.22122111, 0.77877889],
[0.94768725, 0.05231275],
[0.96327122, 0.03672878],
[0.18507726, 0.81492274],
[0.79025684, 0.20974316],
[0.24480067, 0.75519933],
[0.06444165, 0.93555835],
[0.11133741, 0.88866259],
[0.73913651, 0.26086349],
[0.06131543, 0.93868457],
[0.94309103, 0.05690897],
[0.07004774, 0.92995226],
[0.35829366, 0.64170634],
[0.11477531, 0.88522469],
[0.2942615 , 0.7057385],
[0.41079132, 0.58920868],
[0.77690247, 0.22309753],
[0.89545555, 0.10454445],
[0.16130174, 0.83869826],
[0.36873032, 0.63126968],
[0.01685843, 0.98314157],
[0.03340082, 0.96659918],
[0.61623063, 0.38376937],
[0.95189385, 0.04810615],
[0.9643624 , 0.0356376],
[0.23329919, 0.76670081],
[0.07750585, 0.92249415],
[0.21979906, 0.78020094],
[0.86766088, 0.13233912],
[0.77615956, 0.22384044],
[0.05034723, 0.94965277],
[0.48173786, 0.51826214],
[0.13757798, 0.86242202],
[0.73477417, 0.26522583],
[0.9026262 , 0.0973738],


```
[0.98888318, 0.01111682],
[0.02246735, 0.97753265],
[0.94533634, 0.05466366],
[0.01170023, 0.98829977],
[0.01448862, 0.98551138],
[0.06161976, 0.93838024],
[0.02007199, 0.97992801],
[0.40058276, 0.59941724],
[0.73167634, 0.26832366],
[0.92467883, 0.07532117],
[0.16763526, 0.83236474],
[0.39738801, 0.60261199],
[0.71052999, 0.28947001],
[0.04391905, 0.95608095],
[0.71312197, 0.28687803],
[0.3384332 , 0.6615668 ],
[0.89210346, 0.10789654],
[0.46989142, 0.53010858],
[0.65233802, 0.34766198],
[0.15000755, 0.84999245],
[0.27500507, 0.72499493],
[0.9559003 , 0.0440997 ],
[0.13344864, 0.86655136],
[0.31135251, 0.68864749],
[0.93973034, 0.06026966],
[0.33046627, 0.66953373],
[0.93150657, 0.06849343],
[0.91825184, 0.08174816],
[0.94952431, 0.05047569],
[0.03834889, 0.96165111]])
```

```
jovian.commit()
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
confusion_matrix(pred1, val_target, normalize='true')
```

```
array([[0.86486486, 0.13513514],
       [0.10909091, 0.89090909]])
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
tree=DecisionTreeClassifier(random_state=42)
```

```
tree.fit(train_df, train_target)
```

```
DecisionTreeClassifier(random_state=42)
```

```
train_preds1 = tree.predict(train_df)
```

```
tree.score(train_df, train_target)
```


```
1.0
```

```
tree.score(val_df, val_target)
```

```
0.782608695652174
```

```
importance_df = pd.DataFrame({  
    'feature': train_df.columns,  
    'importance': tree.feature_importances_  
}).sort_values('importance', ascending=False)
```

```
plt.title('Feature Importance')  
sns.barplot(data=importance_df.head(10), x='importance', y='feature');
```



```
model = DecisionTreeClassifier(max_depth=8, random_state=42)
```

```
model.fit(train_df, train_target)
```

```
DecisionTreeClassifier(max_depth=8, random_state=42)
```

```
model.score(train_df, train_target)
```

```
0.9539951573849879
```

```
model.score(val_df, val_target)
```

```
0.8152173913043478
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(n_jobs=-1, random_state=42)
```

```
model.fit(train_df, train_target)
```

```
RandomForestClassifier(n_jobs=-1, random_state=42)
```

```
model.score(train_df, train_target)
```


```
1.0
```

```
model.score(val_df, val_target)
```

```
0.9021739130434783
```

```
importance_df = pd.DataFrame({  
    'feature': train_df.columns,  
    'importance': model.feature_importances_  
}).sort_values('importance', ascending=False)
```

```
plt.title('Feature Importance')  
sns.barplot(data=importance_df.head(10), x='importance', y='feature');
```



```
model = RandomForestClassifier(random_state=42, n_jobs=-1, n_estimators=90)
```

```
model.fit(train_df, train_target)
```

```
RandomForestClassifier(n_estimators=90, n_jobs=-1, random_state=42)
```

```
model.score(train_df, train_target)
```

```
1.0
```

```
model.score(val_df, val_target)
```

```
0.9130434782608695
```

```
def test_params(**params):
    model = RandomForestClassifier(random_state=42, n_jobs=-1, n_estimators=90, **params)
    return model.score(train_df, train_target), model.score(val_df, val_target)
```

```
test_params(max_depth=5)
```

```
(0.9043583535108959, 0.8913043478260869)
```

```
test_params(max_depth=26)
```

```
(1.0, 0.9130434782608695)
```

```
test_params(max_features='log2')
```

```
(1.0, 0.9130434782608695)
```

$(1.0, 0.9130434782608695)$

Collecting xgboost

```
|██████████| 173.5 MB 8.9 kB/s eta 0:00:01 |
```

Requirement already satisfied: scipy in /opt/conda/lib/python3.9/site-packages (from xgboost) (1.7.1)

Requirement already satisfied: numpy in /opt/conda/lib/python3.9/site-packages (from xgboost) (1.20.3)

Installing collected packages: xgboost

Successfully installed xgboost-1.5.1

```
model.fit(train_df, train_target)
```

```
/opt/conda/lib/python3.9/site-packages/xgboost/sklearn.py:1224: UserWarning:
```

The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option

use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
[16:29:27] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default
evaluation metric used with the objective 'binary:logistic' was changed from 'error' to
'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```


```
model.score(train_df, train_target)
```

```
model.score(val_df, val_target)
```

0.9130434782608695

```
importance_df = pd.DataFrame({
    'feature': train_df.columns,
    'importance': model.feature_importances_
}).sort_values('importance', ascending=False)
```

```
plt.title('Feature Importance')
sns.barplot(data=importance_df.head(10), x='importance', y='feature');
```



```
import joblib
```

```
heart_model = {
    'model': model,
    'imputer': imputer,
    'scaler': scaler,
    'encoder': encoder,
    'input_cols': input_cols,
    'target_col': target,
    'numeric_cols': numeric_cols,
    'categorical_cols': cat_cols,
    'encoded_cols': encoded_cols
}
```

```
joblib.dump(heart_model, 'heart_model.joblib')
```

['heart_model.joblib']

```
titanic2 = joblib.load('heart_model.joblib')
```

```
jovian.commit()
```

[jovian] Updating notebook "btech60309-19/heart-failure-prediction" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/btech60309-19/heart-failure-prediction>

'<https://jovian.ai/btech60309-19/heart-failure-prediction>'

```
confusion_matrix(model.predict(val_df), val_target, normalize='true')
```

```
array([[0.89473684, 0.10526316],  
       [0.07407407, 0.92592593]])
```

```
jovian.commit()
```